

Essentials of Applied Data Analysis

IPSA-USP Summer School 2018

Sample, Confidence Intervals and Hypothesis
testing in Classical Statistics

Leonardo Sangali Barone
leonardo.barone@usp.br

jan/18

1 Sample and sample distribution of parameters.

CLT and sampling

Central Limit Theorem: if we take enough samples from our population, whose mean of one characteristic of the population (age, income, etc) is μ , and observe the distribution of all of the samples means (\bar{X}), the distribution of the sample means is asymptotically (in the infinite) normal (if, of course, CLT assumptions hold).

Let's take a closer look at the sample and sample distribution of the mean.

Remember: the sample distribution of the mean (of one characteristic of the population) is different from the distribution of the characteristic in the sample

Population *versus* sample

First, let's take a look at the parameters of the population and the sample.




Figure 1: Sampling process

Parameter	Population	Sample estimate
Mean	μ	$\hat{\mu} = \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$
Variance	σ^2	$\hat{\sigma}^2 = s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$

With the sample formulas we can obtain an estimate of the parameters of the population ($\hat{\mu}$ and $\hat{\sigma}^2$ are our best guesses of what the population mean and population variance are).

Population *versus* sample - proportion

In some cases, we are interested in the sample estimation of a proportion, whose notation and formulas can be found below.

Parameter	Population	Sample estimate
Proportion	p	\hat{p}
Variance	$\sigma^2 = p * (1 - p)$	$\hat{\sigma}^2 = s^2 = \hat{p} * (1 - \hat{p})$

Sample distribution of a parameter

Good, we already know the mean and variance of a variable in the population and in a sample. Now, what is the **variance of the sample mean**? Can you see how tricky this question is? Let's call this variance $\sigma_{\bar{x}}^2$ (sigma square of x-bar).

The variance and standard deviation of the sample mean ($\sigma_{\bar{x}}^2$) is simply this:

$$\sigma_{\bar{x}}^2 = \frac{\sigma^2}{n} \text{ and } \sigma_{\bar{x}} = \frac{\sigma}{\sqrt{n}}$$

Where n is the sample size.

Sample distribution of a parameter - HDI

Example: There are 5560 (approx.) municipalities in Brazil. We know (because we have data of the universe of municipalities) that the mean HDI (Human Development Index) of that set of municipalities in 2000 was approximately $\mu_{HDI} = 0.69$ and the standard deviation $\sigma_{IDH} = 0.08$. If we take a random sample of 100 municipalities and observe the sample mean (\bar{x}) what is the standard deviation of the sample mean ($\sigma_{\bar{x}}$)?

$$\sigma_{\bar{x}} = \frac{\sigma_{IDH}}{\sqrt{n}} = \frac{0.08}{\sqrt{100}} = \frac{0.08}{10} = 0.008$$

In our example, we already know what is the population standard deviation, so we can do

$$\sigma_{\bar{x}}^2 = \frac{\sigma^2}{n}$$

What if we don't have the population standard deviation? In large random samples, we can use the sample estimate of the standard deviation as a substitute.

$$\hat{\sigma}_{\bar{x}}^2 = \frac{\hat{\sigma}^2}{n} \text{ or } s_{\bar{x}}^2 = \frac{s^2}{n}$$

Sample distribution of a parameter - proportion

If we are interested in a proportion instead of a mean, then the **variance of the sample proportion \hat{p}** is:

$$\sigma_{\hat{p}}^2 = \frac{\sigma^2}{n} = \frac{\hat{p} * (1 - \hat{p})}{n} \text{ and } \sigma_{\hat{p}} = \frac{\sigma}{\sqrt{n}} = \sqrt{\frac{\hat{p} * (1 - \hat{p})}{n}}$$

where n is the sample size.

Bigger n

On daily research practice, we are often worried about the size of our sample? Why is that?

Because the larger the sample, the smaller the standard deviation of our estimates and the larger is the confidence we got the right number.

Let's see what happens with the distribution of the sample mean in our example of Brazilian municipalities.

Standardization

Before we move to estimation and hypothesis testing, let's learn how to standardize values in the normal distribution.

We know that the normal curve depends on two values: μ and σ , it's mean and standard deviation. To standardize a value means to locate it at one special normal distribution, the Z curve, which is the normal with $\mu = 0$ and $\sigma = 1$

The standardization process of a value x is quite simple:

$$z = \frac{x - \mu}{\sigma}$$

Intuitively, we are "centering" our distribution in the 0 by subtracting μ and making the standard deviation to become a unit by dividing by σ

2 Confidence intervals

Estimation

Point estimate: A single statistic value that is the "best guess" for the parameter value.

Interval estimate: An interval of numbers around the point estimate, that has a fixed "confidence level" of containing the parameter value. Called a confidence interval. (Based on sampling distribution of the point estimate).

Intuitively, confidence intervals are the "range" of values in which the real population parameter is very likely to be (when we don't know the population parameter and we want to learn about it).

Confidence interval

First step:

How to build a confidence interval? First, we take the sample estimate of the parameter. If we want to know the population mean (μ), then we take the population sample (\bar{x}). If we want to know the population proportion

p , then we take the \hat{p} .

Second step:

We define a level of confidence (γ), let's say, 95%. By choosing a level of confidence we are saying that if we sample from our population multiple times, in 95% of the times the population parameter will fall in my interval.

We locate on the Z curve the interval, centered at 0 whose probability is 0.95. This interval is $[-1.96; 1.96]$. We call the limits of the interval critical values (z_γ) .

Final step:

We do the oposite of the standardization process. We multiply this interval we found at the Z curve by the standard deviation of the sample mean ($s_{\bar{x}}$ or $\hat{\sigma}_{\bar{x}}$ and we center it at the sample mean. Our interval with $\gamma = 95\%$ level of confidence is:

$$[\bar{x} - z_\gamma * \hat{\sigma}_{\bar{x}}; \bar{x} + z_\gamma * \hat{\sigma}_{\bar{x}}] = [\bar{x} - 1.96 * \hat{\sigma}_{\bar{x}}; \bar{x} + 1.96 * \hat{\sigma}_{\bar{x}}]$$

Confidence interval - proportion

The confidence interval for proportions looks very similar to the confidence interval for the mean, as long as you do the appropriate substitutions:

$$[\hat{p} - z_\gamma * \hat{\sigma}_{\hat{p}}; \hat{p} + z_\gamma * \hat{\sigma}_{\hat{p}}] = [\hat{p} - 1.96 * \hat{\sigma}_{\hat{p}}; \hat{p} + 1.96 * \hat{\sigma}_{\hat{p}}]$$

Confidence interval

Look at the Z table you have been given. What happens if you choose a greater level of confidence, for example, 99%? What if you choose a lesser level of confidence, for example, 90%?




Figure 2: Confidence Intervals

The confidence interval is itself a random quantity, subject to sampling variability.

In the “long run,” 95% of the confidence intervals for a population mean μ will actually cover μ

Errors and sample size

Look again at the confidence interval.

$$[\bar{x} - z_\gamma * \hat{\sigma}_{\bar{x}}; \bar{x} + z_\gamma * \hat{\sigma}_{\bar{x}}]$$

Let’s define an error ϵ as:

$$\epsilon = z_\gamma * \hat{\sigma}_{\bar{x}}$$

And let’s substitute the standard deviation of the sample mean, $\hat{\sigma}_{\bar{x}}$, by it’s original formula:

$$\epsilon = z_\gamma * \frac{\hat{\sigma}}{\sqrt{n}}$$

Now, imagine that we want to define the maximum error (ϵ) we want and we want to choose a certain level of confidence (which determines z_γ). Since we don’t have any control of the standard deviation of the population (σ), what we have to do is to get a population size (n) that satisfies our restrictions (maximum standard error and level of confidence).

$$\epsilon = z_\gamma * \frac{\sigma}{\sqrt{n}} \Rightarrow n = \frac{\sigma^2 * z_\gamma^2}{\epsilon^2}$$

3 Hypothesis testing

With a good knowledge of how the normal curve is built, sampling e confidence intervals we can now test hypothesis.

We are going to see today only the mean hypothesis testing. If we get the intuition correctly, any other test will be easy.

We also call hypothesis testing as significance tests.

We can generally think in 3 different objectives for hypothesis testing, depending on what we know about the population(s).

- If we know something about the population (mean and variance), we might want to know if a random sample actually comes from that population (and not from some other population);
- If we don't know anything about the population, we might want to learn something (example: mean) or we might want to see if it is equal to our theory/prior belief;
- And finally (and that is more useful in social sciences) we might want to compare two measurements of some the same units (ex: school performance in 2012 vs 2013) or two populations (groups! ex: white-sounding name resumes vs black-sounding) and see if they are equal. In the first case we are assessing the *mean difference* of the 2 measurements. In the second, the *difference of means* of the 2 groups. Tricky, isn't it?

Null Hypothesis

- If we know something about the population (mean and variance), we might want to know if a random sample actually comes from that population (and not from some other population)

$$H_0 : \mu = \text{population mean}$$

- If we don't know anything about the population, we might want to learn something (example: mean) or we might want to see if it is equal to our theory/prior belief. Zero is a common value to be tested.

$$H_0 : \mu = \text{hypothetical value}$$

- Finally (and that is more useful in social sciences) we might want to compare two measurements (X_1 and X_2) of some the same units (mean difference):

$$H_0 : \mu_{(X_1 - X_2)} = 0$$

Or two populations (groups!) and see if they are equal (difference of means).

$$H_0 : \mu_{diff} = \mu_1 - \mu_2 = 0$$

Hypothesis testing - important

- Your null hypothesis (H_0) is the one to be rejected.
- Your null hypothesis is the one that has the = symbol. You can say that it is different by rejecting that it's equal, but not the opposite.
- If you can do it with the sample mean, you can do it with the sample proportions. Just substitute μ for p , \bar{x} for \bar{p} and $\hat{\sigma}_{\bar{x}}$ for $\hat{\sigma}_{\bar{p}}$.
- If you are comparing two measurements or two populations, then your best counterfactual is zero difference. If you are testing dependence, then your best counterfactual is independence. If you are looking for a pattern, then randomness is your counterfactual.

Hypothesis testing - Z or t ?

You will see that for significance tests we will use t-statistics. Remember that for samples over 30 observations it doesn't really matter in practice which one to use. If you don't know which one to use, ask yourself

1. Do you know the population variance? If not, use t .
2. Do you have more than 30 observations? If not, use t .
3. You can use Z if you answered yes to both questions.

Hypothesis testing - mechanics

1. Build your hypothesis
2. Choose a level of significance, $\alpha = 1 - \gamma$, where γ is the level of confidence
3. "Take" your estimated value to the t curve/normal curve by standardizing it. The result will be the t statistics or z statistics:

$$t = \frac{\bar{x} - \mu}{\hat{\sigma}_{\bar{x}}} \text{ or } z = \frac{\bar{x} - \mu}{\sigma_{\bar{x}}}$$

μ , for all purposes, is the value in the null hypothesis.

Hypothesis testing - mechanics for proportions

Substitute μ for p , \bar{x} for \bar{p} and $\hat{\sigma}_{\bar{x}}$ for $\hat{\sigma}_{\bar{p}}$ and you have the proportions test.

$$t = \frac{\hat{p} - p}{\hat{\sigma}_{\hat{p}}} \text{ or } z = \frac{\hat{p} - p}{\sigma_{\hat{p}}}$$

Hypothesis testing - two populations

When we test the difference between the mean of two samples, there are two possible situations. In one, we know the population variance and we also know that both samples come from that population (*equal variances*). Usually, however, we don't know that and we also know that they come from different populations. In those cases, we have to obtain the variance of the difference estimator, $\bar{x}_{diff} = \bar{x}_1 - \bar{x}_2$, which is:

$$\hat{\sigma}_{\bar{x}_{diff}}^2 = \hat{\sigma}_{\bar{x}_1 - \bar{x}_2}^2 = \frac{\hat{\sigma}_{\bar{x}_1}^2}{n_1} + \frac{\hat{\sigma}_{\bar{x}_2}^2}{n_2}$$

The same holds for proportion tests of two populations.